

Harmonogram

9:00 - 10:30 I blok
10:30 - 10:45 Przerwa 15 min
10:45 - 12:00 II blok
12:00 - 12:15 Przerwa 15 min
12:15 - 13:15 III Blok
13:15 - 14:00 Przerwa obiadowa 45m
14:00 - 15:30 IV blok
15:30 - 15:45 Przerwa 15 min
15:45 - 17:00 V blok

Slides: https://docs.google.com/presentation/d/1S-Rs9FOBiDcs5ry3KJa4n7ou6M6mH42jwPPW572hXqU/edit?slide=id.g31bcb11b552_0_208#slide=id.g31bcb11b552_0_208

Stacjonarne:

- Docker i Kubernetes - stacjonarnie:
<https://www.alx.pl/ankiety/7654wvvk>
- Wprowadzenie do Devops - CI/CD, GIT, Jenkins, Ansible, Docker, Kubernetes - stacjonarnie:
<https://www.alx.pl/ankiety/7652efep>

ZDALNIE (ZOOM)

- Wprowadzenie do Devops - CI/CD, GIT, Jenkins, Ansible, Docker, Kubernetes - zdalnie:
<https://www.alx.pl/ankiety/7653mnno>
- Docker i Kubernetes - zdalnie:
<https://www.alx.pl/ankiety/7655heaq>
- Kubernetes w praktyce:
<https://www.alx.pl/ankiety/7656gsiu>

Docker

docker version

hostname

ip a

ls /

cat /etc/os-release

ps aux

```
docker run -it debian bash
```

```
hostname  
ip a  
apt-get update  
apt-get install iproute2
```

```
systemctl status docker.service  
ls -la /var/run/docker.sock  
ls -la /var/run/docker.pid
```

```
pidof dockerd  
cat /var/run/docker.pid  
docker info
```

Docker image registries

<https://hub.docker.com/>
https://hub.docker.com/_/debian

```
docker pull quay.io/prometheus/node-exporter
```

```
docker search ubuntu
```

```
docker search ubuntu --limit 10 --filter=stars=3
```

Docker tags

```
docker pull nginx  
docker pull redis/redis-stack  
docker images
```

```
docker pull redis/redis-stack:7.4.0-v8
```

```
docker pull redis/redis-  
stack@sha256:f3a4ca8891fce481109e663463206d1639a870cba2e5a49a696363abf4e7f95
```

```
docker pull redis/redis-stack:7.2.0-v20
```

docker run

```
docker run --rm -p 8080:80 nginx
```

<http://localhost:8080/>

```
ip addr show dev eth0
```

```
docker run -d --name webserver -p 8081:80 nginx
docker ps
```

<https://github.com/moby/moby/blob/master/internal/namesgenerator/names-generator.go>
<http://localhost:8081/>

```
docker run -d --name redis-rocks -p 8001:8001 redis/redis-stack
```

<http://localhost:8001/>

<http://localhost:8001/redis-stack/workbench>

```
SET training01 k8s-rocks
```

```
GET training01
```

```
docker ps
```

```
docker ps -a
```

```
docker stop redis-rocks
```

```
docker ps -a
```

Container start/stop/restart

```
docker create -p 8080:80 --name nginx007 nginx
docker start containerID
```

```
curl localhost:8080
```

```
docker ps -a
```

```
docker rm fervent_robinson
```

```
docker images
```

```
docker image rm debian:latest
```

```
# remove debian container
```

```
docker rm b6b24e22b38d
```

```
docker image rm debian:latest
```

```
docker container stop 686b7d7722d1  
docker container start 686b7d7722d1
```

```
# running container  
docker rm --force nginx007
```

```
docker exec -it redis-rocks /bin/sh  
redis-cli
```

```
SET Training01 redis-test  
GET Traininig01
```

```
docker exec -it redis-rocks whoami
```

```
docker run --name hello hello-world
```

```
docker logs hello
```

```
docker restart hello  
docker logs hello
```

Container writable-layer

```
docker rm --force redis-rocks
```

```
docker run -d --name redis-rocks -p 8001:8001 redis/redis-stack
```

```
docker volume ls
```

```
docker exec -it redis-rocks ls -R /data/
```

<http://localhost:8001/redis-stack/workbench>
SET mode writable-layer

```
docker restart redis-rocks
```

SET mode writable-layer
SAVE

```
docker exec -it redis-rocks ls -R /data/
```

```
docker restart redis-rocks
```

Bind mount

```
docker run -d --name redis-rocks -p 8001:8001 -v /redis-data/:/data redis/redis-stack
```

```
docker exec -it redis-rocks redis-cli SAVE
```

```
ls /redis-data/
```

Create volume

```
docker volume create redis-vol
```

```
docker volume inspect redis-vol
```

```
docker run -d --name redis-rocks -p 8001:8001 -v redis-vol:/data redis/redis-stack
```

```
docker exec -it redis-rocks redis-cli SAVE
```

```
docker volume inspect redis-vol | jq "[].Mountpoint"
```

```
sudo ls -la /var/lib/docker/volumes/redis-vol/_data
```

```
docker run -d --name redis-rocks --restart unless-stopped -p 8001:8001 -v redis-vol:/data redis/redis-stack
```

```
docker ps
```

```
systemctl restart docker.service
```

```
docker ps
```

```
docker start webserver
```

```
volume-nocopy
```

NFS volume

```
sudo apt-get install -y nfs-kernel-server  
sudo systemctl status nfs-kernel-server.service
```

```
sudo mkdir /nfs
```

```
sudo chown nobody:nogroup /nfs
```

```
ip addr show dev eth0
```

```
sudo vim /etc/exports
```

```
echo "/nfs 10.4.0.0/24(rw)" | sudo tee -a /etc/exports
```

```
cat /etc/exports
```

```
sudo systemctl restart nfs-kernel-server
```

```
nc 10.4.0.XXX 2049 -v
```

```
sudo mkdir /nfs-dump
```

```
sudo mount 10.4.0.10:/nfs /nfs-dump/
```

```
docker volume ls
```

```
docker volume create --driver local \
--opt type=nfs \
--opt o=addr=10.4.0.XXX,rw \
--opt device=/nfs \
redis-nfs
```

```
docker rm --force redis-rocks
```

```
docker run -d --name redis-rocks --restart unless-stopped -p 8001:8001 -v redis-nfs:/data redis/redis-stack
```

```
docker exec -it redis-rocks redis-cli SAVE
```

```
docker volume inspect redis-nfs
```

```
docker inspect redis-rocks
```

```
df -h
```

```
open /nfs
```

Network

```
docker network ls
```

```
docker inspect bridge
```

```
ipcalc 172.17.0.0/16
```

```
docker run --rm -d --name nginx02 nginx
```

```
docker run --rm -it --name curl101 nicolaka/netshoot
```

```
docker inspect bridge
```

```
curl 172.17.0.4

# doesn't resolve
curl nginx02

docker network create demo01

docker run --rm -it --network demo01 --name curl102 nicolaka/netshoot
docker inspect demo01

docker run --rm -d --network demo01 --name nginx03 nginx

# in netshoot container
curl nginx03
```

Dockerfile

```
mkdir -p /home/kurs/src/docker-website
cd /home/kurs/src/docker-website

touch Dockerfile

FROM nginx:1.27.3

LABEL maintainer="nginx@ninja.pl"

ENV NINJA_NGINX_VERSION="0.1"

# comment
RUN apt-get update && \
apt-get install -y \
iproute2 \
wget \
&& rm -rf /var/lib/apt/lists/

COPY index.html /usr/share/nginx/html/index.html

docker build -t nginx-ninja:v1 .

docker images

docker run --rm --name ninja01 -p 8080:80 nginx-ninja:v1
```

```
docker build --build-arg DEST_FILE=ninja.html -t nginx-ninja:v2 .
```

```
docker run --rm --name ninja01 -p 8080:80 nginx-ninja:v2
```

<http://localhost:8080/ninja.html>

Size matters - go app image build

```
git clone https://github.com/max-mulawa/materials.git
cd materials/docker/go-app
```

```
FROM golang:1.24
```

```
LABEL maintainer="random dev"
```

```
RUN apt-get update && \
apt-get install -y curl
```

```
WORKDIR /src
```

```
COPY go.mod go.sum .
RUN go mod download
```

```
COPY cmd/ ./cmd/
COPY pkg/ ./pkg/
```

```
RUN CGO_ENABLED=0 go build -o ./api cmd/web/main.go
```

```
EXPOSE 8080
```

```
ENTRYPOINT ["/src/api"]
docker build -t go-app:v1 .
```

docker images

```
docker run --rm --name goapp -p 8080:8080 go-app:v1
```

<http://localhost:8080/>

```
ps aux | grep api
```

Multi-stage build

```
FROM golang:1.24 AS build
```

```
LABEL maintainer="random dev"
```

```
RUN apt-get update && \  
apt-get install -y curl
```

```
WORKDIR /src
```

```
COPY go.mod go.sum ./  
RUN go mod download
```

```
COPY cmd/ ./cmd/  
COPY pkg/ ./pkg/
```

```
RUN CGO_ENABLED=0 go build -o ./api cmd/web/main.go
```

```
# end result (image build)  
FROM gcr.io/distroless/static-debian12  
COPY --from=build /src/api /app/api
```

```
USER nonroot
```

```
EXPOSE 8080
```

```
ENTRYPOINT ["/app/api"]
```

```
docker build -t go-app:v2 .
```

```
docker pull gcr.io/distroless/static-debian12
```

```
docker run --rm --name goapp -p 8080:8080 go-app:v2
```

```
docker login -u mulawam  
dckr_pat_doQZzGaaECU12OhjSXAtdqS47A
```

```
docker tag go-app:v2 mulawam/go-app:v2-xx  
docker images
```

```
docker push mulawam/go-app:v2-xx
```

```
docker images -f "dangling=true"
```

yaml 101

```
fullName: "Tyson Fury"
```

```
country: "NO"
```

```
active: true
```

```
heavyweightChampion:
```

```
- 2015
```

```
- 2020
```

```
record: [30, 3, 1]
```

```
features:
```

```
  weight: 240
```

```
  height: 204
```

```
fights:
```

```
  - with: Usyk
```

```
    result: loss
```

```
    date: 2024-12-22
```

```
  - with: Usyk
```

```
    result: loss
```

```
    date: 2023-12-05
```

```
notes: >
```

```
  Line 1
```

```
  Line 1 cnt
```

```
notes: |
```

```
  Line 1
```

```
  Line 2
```

```
cat fighter.yaml | yq
```

Docker compose

```
echo -n "mypass" > /tmp/db-password.txt
```

```
services:
```

```
  postgres:
```

```
    image: postgres:16
```

```
    ports:
```

```
      - 5432:5432
```

```
volumes:
  - db-data:/var/lib/postgresql/data
  - ./messages.sql:/docker-entrypoint-initdb.d/messages.sql
secrets:
  - db-password
environment:
  - POSTGRES_PASSWORD_FILE=/run/secrets/db-password
  - POSTGRES_USER=api
api:
  image: go-app:v2
  ports:
    - 8080:8080
  secrets:
    - db-password
  environment:
    - DB_USER=api
    - DB_PASSWORD_FILE=/run/secrets/db-password
    - DB_HOST=postgres
    - DB_PORT=5432
volumes:
  db-data:
secrets:
  db-password:
    file: /tmp/db-password.txt
```

```
docker compose up
curl localhost:8080/messages
```

```
curl -X POST -H 'Content-Type: application/json' -d '{"id":1,"title":"my title", "description":"my body"}'
http://localhost:8080/messages
```

```
curl localhost:8080/messages
```

Published port and nftables

```
docker run --rm -p 8080:80 nginx
ps aux | grep docker-proxy
```

```
sudo nft list ruleset
sudo nft list table ip nat
curl 10.4.0.XX:8080
sudo nft list table ip nat
```

CGroups and memory limits

```
docker run -d --name mem-test polinux/stress stress --vm 1 --vm-bytes 15M
```

```
docker run -d --name mem-test -m 10mb polinux/stress stress --vm 1 --vm-bytes 15M
```

```
docker rm -f mem-test
```

```
docker run -d --name mem-test -m 10mb polinux/stress stress --vm 1 --vm-bytes 8M
```

```
docker stats mem-test
```

```
mount | grep cgrou
```

```
find /sys/fs/cgroup -name *ContainerID*
```

```
pidof stress
```

```
cat /sys/fs/cgroup/system.slice/docker-  
792c1735f2bd973796846c6b10d137fe6719ec777adffb86a7d56a0ff398253d.scope/mem  
ory.max
```

```
cat /sys/fs/cgroup/system.slice/docker-  
792c1735f2bd973796846c6b10d137fe6719ec777adffb86a7d56a0ff398253d.scope/cgroup.procs
```

Kubernetes

Landscape

<https://landscape.cncf.io/>

<https://sonobuoy.io/>

<https://k8s.af/>

<https://kubernetes.io/releases/>

kubectl install

```
mkdir k8s-setup
```

```
cd k8s-setup

KUBECTL_VERSION=v1.33.4
curl -LO https://dl.k8s.io/release/\$KUBECTL\_VERSION/bin/linux/amd64/kubectl

chmod +x kubectl

sudo mv kubectl /usr/local/bin/
kubectl version --client

alias k="kubectl"

source <(kubectl completion bash)
complete -o default -F __start_kubectl k

source ~/.bashrc
k version --client
```

minikube install

```
MINIKUBE_VERSION=v1.37.0
curl -LO https://github.com/kubernetes/minikube/releases/download/\$MINIKUBE\_VERSION/minikube-linux-amd64

chmod +x minikube-linux-amd64
sudo mv minikube-linux-amd64 /usr/local/bin/minikube
minikube status
minikube version

minikube completion bash | sudo tee -a /usr/share/bash-completion/completions/minikube > /dev/null
```

minikube 3 node cluster

```
export K8S_VERSION=1.33.4

minikube start --nodes 3 --memory=2Gb --kubernetes-version=$K8S_VERSION

htop
kubectl get nodes
minikube status

docker ps
docker network ls
```

```
minikube ssh
```

Basic Pod/Deployment/Service

```
mkdir basic  
touch pod.yaml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: basic-nginx  
  labels:  
    training: devops  
spec:  
  containers:  
    - name: c1  
      image: nginx:1.27
```

```
k create -f pod.yaml  
kubectl get pods
```

```
kubectl get pods -o wide
```

```
k port-forward pods/basic-nginx 8080:80  
http://localhost:8080/
```

```
touch deployment.yaml
```

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: basic-nginx-ha  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      training: devops  
      company: pirates  
  template:  
    metadata:  
      labels:  
        training: devops  
        company: pirates  
        ver: "123"
```

```
spec:  
  containers:  
    - name: c1  
      image: nginx
```

```
kubectl create -f deployment.yaml
```

```
kubectl get pods -l company=pirates  
kubectl get deployments.apps
```

```
kubectl get deployment,replicaset
```

```
kubectl get pod --show-labels
```

```
apiVersion: v1  
kind: Service  
metadata:  
  name: basic-nginx  
spec:  
  type: NodePort  
  selector:  
    training: devops  
    company: pirates  
  ports:  
    - port: 8080  
      targetPort: 80  
      nodePort: 30080
```

```
kubectl create -f service.yaml
```

```
minikube ip  
minikube ip -n minikube-m02
```

```
curl 192.168.49.XX:30080
```

```
kubectl logs deployments/basic-nginx-ha --all-pods=true --prefix --since 20s
```

Deployment and rollout strategy

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: infov6  
  annotations:
```

```
kubernetes.io/change-cause: "upgrade to 6.7.1"
spec:
  replicas: 2
  selector:
    matchLabels:
      tier: frontend
      env: dev
  template: # pod definition goes here
    metadata:
      labels:
        tier: frontend
        env: test # broken
    spec:
      containers:
        - name: info
          image: stefanprodan/podinfo:6.6.3
          ports:
            - containerPort: 9898
        - name: net-tools
          image: nicolaka/netshoot
          command: ["/bin/bash"]
          args: ["-c", "while true; do curl localhost:9898/version; sleep 15;done"]
```

```
k create -f deploy.yaml
k get deploy
```

```
k rollout status deployment infov6
k rollout history deployment infov6
k get pod -l tier=frontend
k rollout restart deployment infov6
```

```
k replace -f deploy.yaml
k rollout status deployment infov6
```

Rollback

```
# add non existing podinfo image version
k replace -f deploy.yaml
k get deployments.apps
k get pods
```

```
k describe pod infov6-59b7f8c6bf-9pvv9
```

```
k rollout undo deployment infov6
k get deployments.apps
```

```
k get pod infov6-6db85895bc-7hhmz -o yaml
```

```
image: stefanprodan/podinfo:6.7.1
```

```
k replace -f deploy.yaml
```

Config Maps

```
mkdir cm
```

```
k create configmap podinfo-config --from-literal version=v1 --from-literal msg="Welcome to devops"
```

```
k get cm podinfo-config -o yaml
```

```
touch values.yaml
```

```
database:
```

```
  host: db01.cluster
```

```
  tls: disabled
```

```
touch status.json
```

```
{
  "level": "Warning"
}
```

```
k create configmap podinfo-file-config --from-file val.yaml=values.yaml --from-file status.json
```

```
k get cm
```

```
touch pod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: podinfo-cm
```

```
spec:
```

```
  containers:
```

```
    - name: pi
```

```
      image: stefanprodan/podinfo:6.7.1
```

```
      command:
```

```
        - ./podinfo
```

```
        - --config-path=/configuration
```

```
volumeMounts:
  - name: config-files
    mountPath: /configuration
    readOnly: true
ports:
  - name: http
    containerPort: 9898
    protocol: TCP
env:
  - name: PODINFO_UI_MESSAGE
    valueFrom:
      configMapKeyRef:
        name: podinfo-config
        key: msg
volumes:
  - name: config-files
    configMap:
      name: podinfo-file-config
```

```
k create -f pod.yaml
```

```
k exec -it pod/podinfo-cm -- ls /configuration
```

```
k exec -it pod/podinfo-cm -- cat /configuration/status.json | jq
```

```
k port-forward pod/podinfo-cm 9898:9898
```

<http://localhost:9898/configs>

<http://localhost:9898/>

Namespaces

```
k get namespaces
k get pods --namespace default
k get pods
```

```
k get pods -n kube-system
k get pods -n kube-system -o wide
```

```
k get daemonsets.apps -n kube-system
k get nodes
```

```
k create ns myapi
kubectl get ns
```

```

k get pods -n myapi

k create deployment mydeploy --image stefanprodan/podinfo --replicas 3
k get deployments.apps

k create deployment mydeploy -n myapi --image stefanprodan/podinfo --replicas 3
k get pods -n myapi

k create deployment mydeploy -n myapi --image stefanprodan/podinfo --replicas 3 --dry-run=client

k create deployment mydeploy -n myapi --image stefanprodan/podinfo --replicas 3 --dry-run=client -oyaml

k describe pod mydeploy-c96f9c65d-dv6fv -n myapi

k edit deployments.apps -n myapi mydeploy

```

Services (ClusterIP service)

```

mkdir svc
k create deployment --image nginx --replicas 2 web01

k delete deployments.apps web01
k expose deployment web01 --port 8080 --target-port 80

k run -it --rm pod-tmp --image nicolaka/netshoot -- sh

nslookup web01
curl web01:8080

curl web01.default.svc.cluster.local.:8080

```

Load balancing test (k6)

```

cat <<EOF | tee k6-clusterip.js
import http from 'k6/http';

export const options = {
  vus: 2, //virtual users run concurrently
  iterations: 100, // number of requests throughout the test
  duration: '10s', //duration of the test execution
  noConnectionReuse: true, // disable keep-alive connections
};

```

```
export default function () {
  http.get('http://web01:8080');
}
EOF
```

```
k run -i --rm --image grafana/k6 k6-test-run -- run - < k6-clusterip.js
```

```
k get pods,svc -l app=web01
```

```
k logs --selector app=web01 --prefix --timestamps --since 5m --tail 500 | grep pod/web01-7d78477f84-xvgzk/nginx | wc -l
```

```
k get pod -l app=web01
```

Secrets and private Docker image repository

```
mkdir secrets
cd secrets
```

```
touch deploy.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app-v1
spec:
  selector:
    matchLabels:
      app: app-v1
  template:
    metadata:
      labels:
        app: app-v1
  spec:
    containers:
      - name: c1
        image: mulawam/go-app:v2-mm
    imagePullSecrets:
      - name: docker-reg
```

```
k create -f deploy.yaml
```

```
k get pod -l app=app-v1
```

```
k describe pod app-v1-6f84598564-c65sp | tail
```

```
docker login -u mulawam
dckr_pat_doQZzGaaECUl2OhjlSXAtdqS47A

cat ~/.docker/config.json
echo "bXVsYXdhbTpky2tyX3BhdF9kb1FaekdhYUVDVWwyT2hqbFNYQXRkcVM0N0E=" | base64 -d

kubectl create secret generic docker-reg \
--from-file=.dockerconfigjson=$HOME/.docker/config.json \
--type=kubernetes.io/dockerconfigjson

k get secrets
k get secrets docker-reg -oyaml

k replace -f deploy.yaml
k get pods -l app=app-v1
```

<https://developer.hashicorp.com/vault/tutorials/kubernetes/kubernetes-raft-deployment-guide>
<https://github.com/bitnami-labs/sealed-secrets>

Storage

```
mkdir storage && cd storage
```

Statefulset and emptyDir

```
echo -n root | base64
```

```
touch secret.yaml
```

```
apiVersion: v1
kind: Secret
metadata:
  name: redis-secret
type: Opaque
data:
  password: cm9vdA==
```

```
k create -f secret.yaml
```

```
k get secrets
```

```
touch redis.yaml
```

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: redis
  labels:
    app: redis
spec:
  serviceName: redis
  selector:
    matchLabels:
      service: redis
  template:
    metadata:
      labels:
        service: redis
  spec:
    containers:
      - name: redis
        image: redis/redis-stack:6.2.6-v19
        ports:
          - containerPort: 6379
            name: redis
          - containerPort: 8001
            name: insights
        env:
          - name: REDIS_PASSWORD
            valueFrom:
              secretKeyRef:
                name: redis-secret
                key: password
          - name: REDIS_ARGS
            value: "--requirepass $(REDIS_PASSWORD)"
        volumeMounts:
          - name: redis-data
            mountPath: /data
    volumes:
      - name: redis-data
        emptyDir: {}

---
apiVersion: v1
kind: Service
metadata:
  name: redis
spec:
```

```
ports:  
- port: 6379  
  name: redis  
# this makes its headless and without IP  
clusterIP: None  
selector:  
  service: redis
```

```
k create -f redis.yaml
```

```
k get pods -l service=redis  
k get statefulset
```

```
k exec -it redis-0 -- sh
```

```
redis-cli --user default --pass root  
SAVE
```

```
# other terminal  
k exec redis-0 -- ls /data
```

Hostpath

.....

```
volumes:  
- name: redis-data  
  hostPath:  
    type: DirectoryOrCreate  
    path: /tmp/redis-data
```

```
k replace -f redis.yaml  
k get pods -l service=redis
```

```
k exec redis-0 -- redis-cli --user default --pass root SAVE
```

```
k get pods -l service=redis -owide  
minikube ssh -n minikube-m03
```

```
ls /tmp/redis-data/ -la
```

StorageClass, PV and PVC on minikube

```
k get pods -A  
k get pods -n kube-system  
k get pods -A | grep storage
```

```
k get storageclass
```

```
k get pvc,pv
```

```
k get persistentvolume,persistentvolumeclaims
```

```
touch pvc.yaml
```

```
apiVersion: v1  
kind: PersistentVolumeClaim  
metadata:  
  name: redis-pvc  
spec:  
  accessModes:  
    - ReadWriteOnce  
  storageClassName: standard  
  resources:  
    requests:  
      storage: 1Gi
```

```
k create -f pvc.yaml
```

```
k get pvc,pv
```

<https://kubernetes.io/docs/concepts/storage/storage-classes/>

```
...  
  volumes:  
    - name: redis-data  
    persistentVolumeClaim:  
      claimName: redis-pvc
```

```
k replace -f redis.yaml  
k get pods -l service=redis  
k exec redis-0 -- redis-cli --user default --pass root SAVE
```

```
k describe persistentvolume/pvc-6bdd2b06-b4f9-4e36-b50f-af042fb94c1f
```

```
minikube ssh -n minikube-m03  
k delete -f redis.yaml  
k get pods -l service=redis
```

```
k get pv,pvc
```

Ingress Gateway API

```
mkdir ingress && cd ingress
```

minikube cluster with Calico CNI

```
minikube delete
```

```
docker network rm minikube
```

```
export K8S_VERSION=1.33.4
```

```
minikube start --memory=2Gb --kubernetes-version=$K8S_VERSION --cni=false --network-plugin=cni --extra-config=kubeadm.pod-network-cidr=192.168.0.0/16 --subnet=172.16.0.0/24
```

```
k get pods -A
```

```
export CALICO_VERSION=3.30.2
```

```
k create -f https://raw.githubusercontent.com/projectcalico/calico/v\$CALICO\_VERSION/manifests/operator-crds.yaml
```

```
k api-resources | grep calico
```

```
k create -f https://raw.githubusercontent.com/projectcalico/calico/v\$CALICO\_VERSION/manifests/tigera-operator.yaml
```

```
k get ns
```

```
k get pod -n tigera-operator
```

```
k create -f https://raw.githubusercontent.com/projectcalico/calico/v\$CALICO\_VERSION/manifests/custom-resources.yaml
```

```
watch kubectl get tigerastatus
```

```
k rollout restart deployment -n kube-system coredns
```

```
k get nodes
```

```
minikube node add
```

```
minikube node add
```

```
k get nodes
```

```
touch app.yaml
```

Service (type LoadBalancer) with Metallb

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-multi
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
      tier: frontend
  template:
    metadata:
      labels:
        app: nginx
        tier: frontend
  spec:
    containers:
      - name: ng
        image: nginx
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: lb-svc
spec:
  type: LoadBalancer
  selector:
    app: nginx
    tier: frontend
  ports:
    - targetPort: 80
      port: 8081
```

```
k create -f app.yaml
```

```
k get svc
```

```
minikube addons enable metallb
```

```
k get all -n metallb-system
```

```
172.16.0.0/24
```

```
k get cm -n metallb-system config -o yaml > metallb-config.yaml
```

```
apiVersion: v1
data:
  config: |
    address-pools:
    - name: default
      protocol: layer2
      addresses:
      - 172.16.0.100-172.16.0.120
kind: ConfigMap
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","data":{"config":"address-pools:\n- name: default\n  protocol: layer2\n  addresses:\n  - 172.16.0.100-172.16.0.120"},"kind":"ConfigMap","metadata":{"annotations":{},"name":"config","namespace":"metallb-system"}}
  creationTimestamp: "2025-11-06T15:56:11Z"
  name: config
  namespace: metallb-system
  resourceVersion: "3333"
  uid: b1615b91-c7e8-42d3-8da8-810c07aa0cfe
```

```
k replace -f metallb-config.yaml
```

```
k rollout restart deployment -n metallb-system controller
```

```
curl 172.16.0.100:8081
```

Calico Gateway API (Ingress)

```
minikube status
minikube start
```

```
k get nodes
```

```
k rollout restart daemonset -n calico-system calico-node
```

```
k get pods -A -owide
```

```
minikube node delete minikube-m0X
```

```
k get pods -A -owide | grep m03
```

```
touch tigera-gateway.yaml
```

```
apiVersion: operator.tigera.io/v1
kind: GatewayAPI
metadata:
  name: tigera-secure
```

```
k create -f tigera-gateway.yaml
```

```
watch kubectl get tigerastatus
```

```
k get ns
```

```
k get all -n tigera-gateway
```

```
k get pod -A -owide
```

```
touch podinfo-gateway.yaml
```

```
apiVersion: gateway.networking.k8s.io/v1
kind: Gateway
metadata:
  name: podinfo-gateway
  namespace: default
spec:
  gatewayClassName: tigera-gateway-class
  listeners:
    - name: http
      port: 80
      protocol: HTTP
      hostname: podinfo.public
    - name: https
      port: 443
      protocol: HTTPS
      hostname: podinfo.public
      tls:
        mode: Terminate
        certificateRefs:
          - name: podinfo-public-tls
            kind: Secret
```

```
sudo apt install mkcert libnss3-tools -y
mkcert --install
```

```
mkcert podinfo.public
```

```
k create secret tls podinfo-public-tls --cert ./podinfo.public.pem --key podinfo.public-key.pem  
k get secrets
```

```
k create -f podinfo-gateway.yaml
```

```
k get all -n tigera-gateway
```

```
k replace --force -f metallb-config.yaml
```

```
touch podinfo-app.yaml
```

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: pod-info-public  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      component: pod  
      app: info-public  
  template:  
    metadata:  
      labels:  
        component: pod  
        app: info-public  
    spec:  
      containers:  
        - name: c1  
          image: stefanprodan/podinfo
```

```
---  
apiVersion: v1  
kind: Service  
metadata:  
  name: podinfo-public-svc  
spec:  
  selector:  
    component: pod  
    app: info-public  
  ports:  
    - name: http  
      port: 8080  
      targetPort: 9898
```

```
k create -f podinfo-app.yaml
```

```
k get pods
```

```
touch http-route.yaml
```

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: podinfo-http-route
  namespace: default
spec:
  parentRefs:
    - name: podinfo-gateway
      namespace: default
      port: 80
      sectionName: http
  hostnames:
    - podinfo.public
  rules:
    - matches:
        - path:
            type: PathPrefix
            value: /
  backendRefs:
    - name: podinfo-public-svc
      port: 8080
```

```
k get pod -n metallb-system
```

```
k get gateway
```

```
curl -v -H "Host: podinfo.public" http://172.16.0.101
```

```
touch https-route.yaml
```

```
apiVersion: gateway.networking.k8s.io/v1
kind: HTTPRoute
metadata:
  name: podinfo-route
  namespace: default
spec:
  parentRefs:
    - name: podinfo-gateway
      namespace: default
      port: 443
      sectionName: https
```

```
hostnames:  
- podinfo.public  
rules:  
- matches:  
  - path:  
    type: PathPrefix  
    value: /  
backendRefs:  
- name: podinfo-public-svc  
  port: 8080
```

```
kubectl create -f https-route.yaml
```

```
echo "172.16.0.101 podinfo.public" | sudo tee -a /etc/hosts
```

```
curl https://podinfo.public
```

```
kubectl get httproute  
kubectl get pods -n tigera-gateway
```

```
curl https://podinfo.public
```

```
https://kind.sigs.k8s.io/  
https://github.com/kelseyhightower/kubernetes-the-hard-way
```

```
https://github.com/kubernetes-sigs/kubespray
```

```
export CP="10.4.0.XXX" # control plane  
export W1="10.4.0.XXX" # worker 1  
export W2="10.4.0.XXX" # worker 2
```

```
source ~/.bashrc
```

```
echo $CP $W1 $W2
```

```
cd
```

```
ssh-keygen
```

```
for ip in $CP $W1 $W2; do  
  ssh-copy-id -o StrictHostKeyChecking=accept-new $ip  
done
```

```
for ip in $CP $W1 $W2; do  
  ssh ${ip} hostname
```

```
ssh ${ip} ip addr show dev eth0
done

ssh $CP

sudo ufw status

swapon --show
sudo sysctl net.ipv4.ip_forward
sudo sysctl net.ipv4.ip_forward=1
sudo sysctl net.ipv4.ip_forward

echo "net.ipv4.ip_forward = 1" | sudo tee /etc/sysctl.d/k8s.conf

sudo modprobe overlay

cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
EOF

sudo apt-get update

sudo install -m 0755 -d /etc/apt/keyrings
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc
sudo chmod a+r /etc/apt/keyrings/docker.asc

{
  echo \
    "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc] \
https://download.docker.com/linux/ubuntu \
    $(./etc/os-release && echo "$VERSION_CODENAME") stable" | \
  sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
}
sudo apt-get update

sudo apt-get -y install containerd.io

containerd config default

containerd config default | sudo tee /etc/containerd/config.toml

cat /etc/containerd/config.toml

sudo sed -i "s/SystemdCgroup = false/SystemdCgroup = true/g" /etc/containerd/config.toml

cat /etc/containerd/config.toml | grep Cgroup
```

```
sudo systemctl restart containerd
sudo systemctl status containerd

export K8S_MINOR_VERSION=1.33
curl -fsSL https://pkgs.k8s.io/core:/stable:/v\$K8S\_MINOR\_VERSION/deb/Release.key | sudo gpg --dearmor
-o /etc/apt/keyrings/kubernetes-apt-keyring.gpg

echo "deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:/v\${K8S\_MINOR\_VERSION}/deb//" | sudo tee /etc/apt/sources.list.d/kubernetes.list

https://semver.org/

sudo apt-get update

sudo apt-cache madison kubeadm

export K8S_VERSION=1.33.4
sudo apt-get install -y kubelet=$K8S_VERSION-1.1 kubeadm=$K8S_VERSION-1.1
kubectl=$K8S_VERSION-1.1

sudo apt-mark hold kubelet kubeadm kubectl

sudo kubeadm init phase preflight

sudo sed -i "s/pause:3.8/pause:3.10/g" /etc/containerd/config.toml

sudo systemctl restart containerd
sudo systemctl status containerd

cat <<EOF | tee kubeadm.config
apiVersion: kubeadm.k8s.io/v1beta4
kind: ClusterConfiguration
kubernetesVersion: v1.33.4
apiServer:
  extraArgs:
    - name: advertise-address
      value: ${CP}
networking:
  podSubnet: 192.168.0.0/16
  serviceSubnet: 10.96.0.0/12
---
apiVersion: kubeproxy.config.k8s.io/v1alpha1
kind: KubeProxyConfiguration
mode: nftables
EOF
```

```
rm kubeadm.config
exit

hostname
echo $CP

scp ./kubeadm.config $CP:/home/kurs
ssh $CP ls /home/kurs

ssh $CP

sudo kubeadm init --config kubeadm.config --dry-run | more

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

kubectl get nodes

kubectl describe node 00control

kubectl get pod -A -o wide

scp $CP:/home/kurs/.kube/config $HOME/.kube/config

cat ~/.kube/config

k get nodes

mkdir setup-calico && cd setup-calico

export CALICO_VERSION=3.30.2
kubectl create -f https://raw.githubusercontent.com/projectcalico/calico/v\$CALICO\_VERSION/manifests/operator-crds.yaml
kubectl create -f https://raw.githubusercontent.com/projectcalico/calico/v\$CALICO\_VERSION/manifests/tigera-operator.yaml

k get -n tigera-operator all

wget https://raw.githubusercontent.com/projectcalico/calico/v\$CALICO\_VERSION/manifests/custom-resources.yaml
cat custom-resources.yaml

sed -i "s/VXLANCrossSubnet/VXLAN/g" ./custom-resources.yaml
```

```
sed -i '/calicoNetwork:/a\\ \\ \\ linuxDataplane: Nftables' ./custom-resources.yaml
```

```
kubectl create -f custom-resources.yaml
```

```
kubectl get nodes  
kubectl get tigerastatus  
kubectl get nodes
```

Worker nodes setup

```
tee depend.sh > /dev/null <<'DEPS'
```

```
#!/bin/bash  
set -euo pipefail
```

```
export DEBIAN_FRONTEND=noninteractive
```

```
sudo sysctl net.ipv4.ip_forward=1
```

```
cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf  
net.ipv4.ip_forward = 1  
EOF
```

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf  
overlay  
EOF
```

```
sudo modprobe overlay
```

```
sudo apt-get update  
sudo install -m 0755 -d /etc/apt/keyrings  
sudo curl -fsSL https://download.docker.com/linux/ubuntu/gpg -o /etc/apt/keyrings/docker.asc  
sudo chmod a+r /etc/apt/keyrings/docker.asc
```

```
{  
    echo \  
    "deb [arch=$(dpkg --print-architecture) signed-by=/etc/apt/keyrings/docker.asc]  
https://download.docker.com/linux/ubuntu \  
    $(./etc/os-release && echo "$VERSION_CODENAME") stable" | \  
    sudo tee /etc/apt/sources.list.d/docker.list > /dev/null  
}
```

```
sudo apt-get update  
sudo apt-get -y install containerd.io
```

```
sudo mkdir -p /etc/containerd
containerd config default | sudo tee /etc/containerd/config.toml

sudo sed -i "s/SystemdCgroup = false/SystemdCgroup = true/g" /etc/containerd/config.toml
sudo sed -i "s/pause:3.8/pause:3.10/g" /etc/containerd/config.toml

sudo systemctl restart containerd

export K8S_MINOR_VERSION=1.33 # K8S_MINOR_VERSION tag

curl -fsSL https://pkgs.k8s.io/core:/stable:v\$K8S\_MINOR\_VERSION/deb/Release.key | sudo gpg --dearmor
--yes --batch -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg

echo "deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg] https://pkgs.k8s.io/core:/stable:v\${K8S\_MINOR\_VERSION}/deb//" | sudo tee /etc/apt/sources.list.d/kubernetes.list

sudo apt-get update

export K8S_VERSION=1.33.4 # K8S_VERSION tag
sudo apt-get install -y kubelet=$K8S_VERSION-1.1 kubeadm=$K8S_VERSION-1.1
kubectl=$K8S_VERSION-1.1
sudo apt-mark hold kubelet kubeadm kubectl

kubectl version --client=true
kubeadm version
kubelet --version
DEPS

for host in $W1 $W2; do
    scp ./depend.sh $host:./depend.sh | tee -a deps-install.log
    echo "Running depend.sh on $host ..." | tee -a deps-install.log
    ssh $host ./depend.sh | tee -a deps-install.log
done

# check the log output
cat deps-install.log | grep -E "(kubeadm|kubelet|10\.)"

ssh $W1 kubeadm version
ssh $W2 kubeadm version

ssh $CP sudo kubeadm token create --print-join-command

JOIN_CMD=$(ssh $CP sudo kubeadm token create --print-join-command)
echo $JOIN_CMD
```

```
for worker in $W1 $W2; do
  echo "joining $worker node to k8s cluster" | tee -a join-cluster.log
  ssh $worker sudo $JOIN_CMD | tee -a join-cluster.log
done
```

k get nodes

```
sudo apt-get install curl gpg apt-transport-https --yes
curl -fsSL https://packages.buildkite.com/helm-linux/helm-debian/gpgkey | gpg --dearmor | sudo tee
/usr/share/keyrings/helm.gpg > /dev/null
echo "deb [signed-by=/usr/share/keyrings/helm.gpg] https://packages.buildkite.com/helm-linux/helm-debian/any/ any main" | sudo tee /etc/apt/sources.list.d/helm-stable-debian.list
sudo apt-get update
sudo apt-get install helm
```

helm version

<https://k9scli.io/>

Kubernetes Dashboard

helm repo add kubernetes-dashboard <https://kubernetes.github.io/dashboard/>

```
helm install dashboard kubernetes-dashboard/kubernetes-dashboard --create-namespace --namespace
kubernetes-dashboard
```

helm list -A

k get all -n kubernetes-dashboard

k port-forward -n kubernetes-dashboard svc/dashboard-kong-proxy 8443:443

<https://localhost:8443/>

```
k get serviceaccounts -n tigera-operator
k get clusterrole | grep admin
```

touch rbac.yaml

```
# Service Account with the name admin-user in namespace kubernetes-dashboard first
apiVersion: v1
kind: ServiceAccount
metadata:
```

```
name: admin-user
namespace: kubernetes-dashboard
---
# In most cases after provisioning the cluster using kops, kubeadm or any other popular tool, the ClusterRole
cluster-admin already exists in the cluster.
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: admin-user
roleRef:
  apiGroup: rbac.authorization.k8s.io
  kind: ClusterRole
  name: cluster-admin
subjects:
- kind: ServiceAccount
  name: admin-user
  namespace: kubernetes-dashboard
```

```
k create -f rbac.yaml
```

k get sa -n kubernetes-dashboard

```
k create -n kubernetes-dashboard token admin-user
```

```
fg
ctrl+C
```

Dynamic storage (NFS server)

NFS Server on jumpbox

```
sudo apt install -y nfs-kernel-server
sudo systemctl status nfs-kernel-server
sudo mkdir -p /nfs
sudo chown nobody:nogroup /nfs

echo "/nfs 10.4.0.0/24(rw)" | sudo tee -a /etc/exports
sudo systemctl restart nfs-kernel-server
```

NFS client on workers

```
for host in $CP $W1 $W2; do
  echo "Installing NFS client sh on $host ..." | tee -a nfs-install.log
  ssh $host sudo apt-get install -y nfs-common | tee -a nfs-install.log
done
```

NFS Provider installed through helm chart

<https://github.com/kubernetes-sigs/nfs-subdir-external-provisioner/blob/master/charts/nfs-subdir-external-provisioner/README.md>

```
ip addr show dev eth0
```

<https://github.com/kubernetes-sigs/nfs-subdir-external-provisioner/blob/master/charts/nfs-subdir-external-provisioner/values.yaml#L10C1-L12C21>

```
helm repo add nfs-subdir-external-provisioner https://kubernetes-sigs.github.io/nfs-subdir-external-provisioner
```

```
helm install nfs-subdir-external-provisioner nfs-subdir-external-provisioner/nfs-subdir-external-provisioner \
--create-namespace \
--namespace nfs-provisioner \
--set nfs.server=10.4.0.XXX \
--set nfs.path=/nfs/
```

```
helm list -A
```

```
k rollout restart deployment nfs-subdir-external-provisioner -n nfs-provisioner
```

```
k get storageclass
```

```
kubectl patch storageclass nfs-client -p '{"metadata": {"annotations": {"storageclass.kubernetes.io/is-default-class": "true"}}}'
```

```
k get storageclass
```

```
k get all -n nfs-provisioner
```

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: redis
  labels:
    app: redis
spec:
  serviceName: redis
```

```
selector:
matchLabels:
  service: redis
template:
  metadata:
    labels:
      service: redis
spec:
  containers:
    - name: redis
      image: redis/redis-stack:6.2.6-v19
      ports:
        - containerPort: 6379
          name: redis
        - containerPort: 8001
          name: insights
      env:
        - name: REDIS_PASSWORD
          valueFrom:
            secretKeyRef:
              name: redis-secret
              key: password
        - name: REDIS_ARGS
          value: "--requirepass $(REDIS_PASSWORD)"
      volumeMounts:
        - name: redis-data
          mountPath: /data
    volumes:
      - name: redis-data
volumeClaimTemplates:
  - metadata:
      name: redis-data
    spec:
      accessModes: [ "ReadWriteOnce" ]
      storageClassName: "nfs-client"
      resources:
        requests:
          storage: 1Gi
---  
apiVersion: v1
kind: Service
metadata:
  name: redis
spec:
  ports:
    - port: 6379
```

```
name: redis
# this makes its headless and without IP
clusterIP: None
selector:
  service: redis
```

```
kubectl create -f secret.yaml -f redis.yaml
```

```
kubectl get pvc,pv
```

```
kubectl exec redis-0 -- redis-cli --user default --pass root SAVE
```

```
tree /nfs
```

<https://github.com/gimlet-io/onechart>

Network Policies

CNI

<https://github.com/flannel-io/flannel>

<https://cilium.io/use-cases/cni/>

Policies

```
mkdir policies
```

```
cd policies
```

```
touch manifest.yaml
```

```
apiVersion: v1
kind: Namespace
metadata:
  name: experiments
  labels:
    experiment: v1
---
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
name: pod-info
namespace: experiments
spec:
  replicas: 2
  selector:
    matchLabels:
      component: pod
      app: info
  template:
    metadata:
      labels:
        component: pod
        app: info
    spec:
      containers:
        - name: c1
          image: stefanprodan/podinfo
```

```
---
```

```
apiVersion: v1
kind: Service
metadata:
  name: podinfo-svc
  namespace: experiments
spec:
  selector:
    component: pod
    app: info
  ports:
    - name: http
      port: 9898
      targetPort: 9898
```

```
kubectl create -f manifest.yaml
```

```
kubectl get all -n experiments
```

```
kubectl run -n default -it --rm test-default --image nicolaka/netshoot
```

```
curl podinfo-svc
curl podinfo-svc.experiments:9898
```

```
cat /etc/resolv.conf
nslookup podinfo-svc.experiments
```

```
curl podinfo-svc.experiments.svc.cluster.local.:9898
```

```
k run -it --rm test-exper -n experiments --image nicolaka/netshoot -- sh  
curl podinfo-svc:9898
```

```
touch ingress-policy.yaml
```

```
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: only-experiments  
  namespace: experiments  
spec:  
  podSelector:  
    matchLabels:  
      app: info  
      component: pod  
  policyTypes:  
    - Ingress
```

```
k get networkpolicies -n experiments
```

```
k describe networkpolicies/only-experiments -n experiments
```

```
k -n calico-system port-forward svc/whisker 8081:8081 &
```

<http://localhost:8081/flow-logs/denied-flows>

```
curl --connect-timeout 2s podinfo-svc:9898
```

```
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: only-experiments  
  namespace: experiments  
spec:  
  podSelector:  
    matchLabels:  
      app: info  
      component: pod  
  policyTypes:  
    - Ingress  
ingress:  
  - from:  
    - namespaceSelector:  
      matchLabels:  
        experiment: default
```

```
k label namespaces default experiment=default  
k get ns default --show-labels
```

```
k replace -f ingress-policy.yaml
```

```
curl podinfo-svc.experiments.svc.cluster.local.:9898  
curl --connect-timeout 2s podinfo-svc:9898
```

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: only-experiments
```

```
  namespace: experiments
```

```
spec:
```

```
  podSelector:
```

```
    matchLabels:
```

```
      app: info
```

```
      component: pod
```

```
  policyTypes:
```

```
    - Ingress
```

```
  ingress:
```

```
    - from:
```

```
      - namespaceSelector:
```

```
        matchExpressions:
```

```
          - key: experiment
```

```
            operator: In
```

```
            values: ["default","v1"]
```

```
    # matchLabels:
```

```
    # experiment: default
```

```
curl podinfo-svc.experiments.svc.cluster.local.:9898
```

```
curl --connect-timeout 2s podinfo-svc:9898
```

<https://docs.tigera.io/calico/latest/network-policy/encrypt-cluster-pod-traffic>